

EXACT 2026 Data Disclosure Statement

Team: AI WITH BRO | **Scope:** Training data, retrieval corpora, preprocessing inputs, and evaluation assets

This statement discloses all external data sources used by the EXACT-2026 pipeline. The system does not use GPT, Claude, Gemini, commercial LLM APIs, or any closed-source model for training, data generation, preprocessing, retrieval, evaluation, or inference. "OpenAI-compatible" in the codebase refers only to the local HTTP serving protocol used by llama-server.

Source Register

Source	Role in pipeline	Scope used	Processing and controls	Disclosure status
Official EXACT 2026 Dataset	Primary task data for Type 1 logic and Type 2 physics; validation of API schema and answer format.	Logic: 411 records / 808 questions. Physics: organizer-provided text-only CSV.	Parsed deterministically from JSON/CSV. Physics rows whose IDs start with QA are filtered out according to the official Q&A ruling.	Fully disclosed as organizer-provided data.
FOLIO Dataset	Improves natural-language-to-first-order-logic mapping and Z3 entailment patterns for Type 1 logic.	Approx. 1,082 records after availability and preprocessing filters.	Premises and conclusions are converted into formal logic examples using deterministic Python scripts and open-source symbolic tooling.	Public source: https://huggingface.co/datasets/yale-nlp/FOLIO
PhysicsFormulae Repository	Static formula and constants reference for Type 2 physics retrieval.	Approx. 28 verified formula/constant records.	Extracted, normalized, verified, and converted into RAG JSONL records.	Public source: https://github.com/BenjaminTMilnes/PhysicsFormulae
Electro Textbook Dataset	Supplementary physics examples for SymPy code-generation and explanation training.	242 physics problems with solutions and verified SymPy code.	Internally digitized and normalized into supervised examples. Must be excluded unless a complete bibliography is attached.	Conditional disclosure required before final submission.
RAG Knowledge Base	Runtime retrieval for physics formalization.	Collections: physics_formulas, physics_examples.	Built from official physics data, PhysicsFormulae, and optionally Electro only if fully disclosed. Uses embeddings/reranking; no closed-source LLM rewriting.	Derived corpus; sources listed above.

Closed-Source LLM Attestation

No closed-source or commercial LLM is used at any point in the pipeline. This includes: generating synthetic training data, rewriting datasets, preprocessing examples, building RAG content, running evaluation, or serving predictions. Runtime inference is performed only with self-hosted open-source 7B-class models through llama.cpp / llama-server.

Submission Risk Note

The Electro textbook dataset is the only conditional source. If the team cannot provide title, author or publisher, edition/year, URL or acquisition source, and page/section range for every included item, Electro-derived artifacts should be removed from fine-tuning, retrieval indexes, and the final live API package.