

EXACT 2026 Data Disclosure Document

Team: AI WITH BRO | **Scope:** training, fine-tuning, retrieval, evaluation, and deployment data

Complete Data Registry

Dataset / model source	Size used	Purpose	Source and processing
EXACT 2026 official release (May 15)	Type 1: 411 records / 808 questions; Type 2: 1,352 problems	Fine-tuning, validation, format definition, exact-match retrieval for released Type 1 examples	ura.hcmut.edu.vn/exact . Parsed locally; the updated release already removes invalid Type 2 QA rows.
FOLIO	1,204 train/validation examples loaded; 199 executable examples retained for coder tuning	Logic NL-to-FOL/Z3 and explanation tuning	huggingface.co/datasets/yale-nlp/FOLIO . Converted locally to executable Z3 scripts; invalid conversions are discarded.
Internal electro textbook-derived corpus	242 records	Additional Physics/SymPy and explanation tuning	Manually collected text problems. Per-record source URLs were not retained; this provenance limitation is disclosed. Converted locally to SymPy templates and verified before use.
Fine-tuned Qwen GGUF models	Qwen2.5-Coder-7B and Qwen2.5-7B-Instruct, Q4_K_M	Runtime code generation and Logic MCQ answering	Public weights: huggingface.co/HoangKhangHCMUS/exact-2026 . Self-hosted with llama-server.

Generated Artifacts And Runtime Retrieval

- Fine-tuning outputs contain 1,546 coder samples and 2,798 instruct samples. Scripts and aggregate statistics are included in the repository.
- No closed-source model generated synthetic data, and no larger teacher-model distillation was used.
- No external RAG corpus or persisted vector database is active. Runtime retrieval is limited to transparent exact full-input matching over the disclosed released Type 1 corpus; unseen inputs continue through the reasoning pipeline.
- Runtime inference does not call proprietary LLM APIs or third-party inference providers. Z3, SymPy, deterministic formula rules, and the two self-hosted Qwen models are used.

Controls

All preprocessing runs locally. Generated Python is AST-validated, executed in an isolated subprocess with time/output/resource limits, and covered by automated tests. Model metadata is auditable through `/v1/models`.